FAQ

- **Is this real user data?**

No, everything is synthetic. User IDs, item IDs, timestamps, and click events are generated by a simulator; nothing here is tied to real people or real products. Treat it purely as a statistical modeling problem.

- **Are AI coding tools (Copilot, ChatGPT, Claude Code, etc.) allowed?**

Yes, of course. Use whatever tools help you build and debug your model.

- **What am I actually predicting, and how is it scored?**

For each row of evaluation.csv (one candidate item shown to a user in one event), predict the probability it was clicked. The primary metric is global ROC AUC across all evaluation rows. A constant prediction scores ~0.5 AUC by construction.

- **Is the leaderboard score I see during the competition final?**

No, it's a provisional score on a subset of the evaluation set. The final ranking is computed on the full evaluation set (including rows held back from the live leaderboard) after the competition closes, so don't over-tune to the live number.

- **What format does my submission need to be in?**

A CSV with header rowid,predicted_probability, exactly one row per rowid in evaluation.csv, each exactly once, probabilities as finite numbers in [0, 1]. Run validate_submission.py before submitting — malformed, missing, duplicate, or out-of-range submissions are rejected, not auto-repaired.

- **Can I use rowid or event_id as a feature?**

Not recommended. Both are random identifiers assigned at export time and carry no information. Don't waste time trying to decode them.

- **How many submissions am I allowed?**

10 submissions per day, per team.